

EB-H-JEPA: Energy-Based Hamiltonian JEPA World Models

A Failure Atlas and a Controlled Sample-Efficiency Test in Crafter

Sehaj Randhir Singh
NYU Tandon School of Engineering
sehaj.randhir.singh@nyu.edu

Abstract

World models that operate in a learned latent space (JEPA; LeCun 2022) are the leading bet for sample-efficient, long-horizon model-based RL (DreamerV3, Hafner et al. 2023). A recurring hypothesis is that imposing physical structure on the latent predictor — Hamiltonian, metriplectic, or GENERIC — should make learned dynamics faithful and stable. We test this hypothesis twice, with instruments of increasing fidelity. First, on a nonlinearly lifted Lorenz-63 system, we audit five structurally distinct latent predictors under a chaos-aware protocol. The result is a failure atlas: each constraint class fails, but on a property disjoint from the one it enforces. A rigid Hamiltonian predictor freezes and cannot dissipate; a naive metriplectic predictor sits on a dead $LL^\top$ saddle with dissipation pinned at exactly zero; fixing both bugs (constant divergence-free J , nonzero R initialization) makes dissipation engage and recovers the true phase-space contraction -13.667 to $\sim 1\%$ accuracy — yet spectral alignment still cannot pin individual Lyapunov exponents. Second, and constructively, we insert the fixed-metriplectic predictor into a compact DreamerV3-style agent and measure sample efficiency on Crafter (Hafner 2021), the benchmark of the world-model group that introduced RSSM. If structure helps, this is the first positive evidence that physical inductive biases pay off in an RL world model; if it hurts, that is the honest negative the field needs. All code, runs, and evidence are public; every claim traces to a JSON artifact.

1 Introduction

JEPA-style world models [12, 2, 5] and their model-based RL instantiations [9, 10, 3] share a core bet: a latent predictor learned from observations can model dynamics that stay stable and meaningful over long horizons. The predictor is the crux. DreamerV3’s categorical RSSM learns it from data alone; a growing line of work asks whether injecting physical structure — symplectic/Hamiltonian form [6, 4, 16], metriplectic/GENERIC dissipation [15, 17, 11], or learned divergence gates [7] — yields predictors that are more faithful and more sample-efficient.

This paper reports a two-part test of that hypothesis, built to be reproducible end to end.

Part 1: a failure atlas on a known chaotic system. We train five structurally distinct latent predictors on Lorenz-63 lifted into a 64-D latent chart and audit each under a chaos-aware protocol: rollout boundedness, leading Lyapunov exponent, phase-space contraction, attractor overlap, and dissipation engagement. The honest headline is negative, and more precise than “structure doesn’t help”: *every constraint class fails, but on a property disjoint from the one it enforces*. A rigid Hamiltonian predictor (fixed canonical J , learned H) is conservative by construction and cannot dissipate; it freezes into a quasi-periodic tube with near-zero motion. A naive metriplectic predictor learns R through $R = LL^\top$ but sits on the dead saddle where $L = 0$: R stays exactly 0, dissipation never engages. Fixing both bugs — constant divergence-free J (closing the learned-skew loophole) and nonzero L initialization (breaking the dead saddle) — makes R engage and recovers the true contraction -13.667 to $\sim 1\%$. Yet even the best arm cannot pin individual Lyapunov exponents: differentiable QR spectral shaping controls only a finite-horizon statistic, not the asymptotic Oseledets spectrum [14]. And overlap metrics reward dead models: a frozen predictor sits inside the encoded attractor and posts the best Chamfer score while being dynamically inert.

Part 2: a controlled sample-efficiency test in Crafter. A failure atlas on Lorenz tells us where structure fails in isolation; it does not tell us whether structure helps or hurts an *actual RL world model*, where the predictor co-adapts with an actor trained by latent imagination. We build a compact but faithful reimplement of the DreamerV3 recipe — CNN encoder/decoder, recurrent memory, categorical posterior, KL free bits, symlog targets, TD(λ) imagination — in which the stochastic prior is pluggable. The baseline is DreamerV3’s exact categorical RSSM; the treatment is a continuous latent whose prior mean advances under the fixed-metriplectic map from Part 1. Everything else is identical. We measure sample efficiency on Crafter [8], the open-world benchmark introduced by the RSSM group itself. The result — whether structure helps, hurts, or is neutral — is a controlled answer to the field’s open question, and either way it is publishable. Preliminary CPU verification (Sec. 4.2) confirms both arms train and the pipeline is sound; the headline numbers (Table 2) come from the 1-hour T4 runs whose exact commands are shipped in the repo.

2 Background

JEPA. Joint-Embedding Predictive Architectures [12] predict in representation space, not pixel space: an encoder produces a latent z_t from observation x_t , and a predictor advances it. The EB-JEPA line [2, 5] shows strong representation learning without reconstruction.

DreamerV3. Hafner et al. [9] train a recurrent state-space model (RSSM) — categorical stochastic prior $p(z_t | h_t)$, posterior $q(z_t | h_t, x_t)$, GRU memory h — jointly with reconstruction, reward, and continue heads, then train an actor-critic by “imagining” rollouts in the learned latent. It is the sample-efficiency reference for visual control [8].

Metriplectic / GENERIC. Dissipative physical systems are not Hamiltonian; they follow the metriplectic (GENERIC) form $\dot{z} = J\nabla H - R\nabla S$ [15], with J antisymmetric (energy-conserving flow), R symmetric positive semidefinite (dissipation), and H, S the energy and entropy potentials. Learned versions [17, 11, 7] parameterize J and R ; the two

failure modes isolated in this paper — learned J exploiting the skew loophole, and $R = LL^\top$ sitting at the dead $L = 0$ saddle — are exactly the traps a naive parameterization falls into.

3 Method

3.1 Latent predictors (Part 1)

All arms share: a 64-D latent chart (nonlinear diffeomorphic lift of Lorenz-63), an MLP dynamics predictor, SIGReg anti-collapse regularization, and identical training budgets. They differ only in the algebraic class of the predictor:

B Unconstrained MLP $z' = f_\theta(z)$.

C Rigid Hamiltonian $z' = z + dt J \nabla_z H(z, h)$, fixed canonical J ; conservative, cannot dissipate.

D Naive metriplectic, learned J , $R = L(z)L(z)^\top$, L initialized at zero.

E *Fixed* metriplectic: constant divergence-free J ($\text{div}(J \nabla H) = 0$ exactly), L initialized at 0.05, learned H, S conditioned on GRU state h .

F E plus differentiable-QR spectral alignment on the rollout Jacobian.

Chaos-aware evaluation: rollout boundedness, motion ratio $m = \mathbb{E}|\Delta z|/\mathbb{E}|z|$, leading exponent λ_1 by the QR method, contraction $\text{tr}(R)/n$ vs. true divergence -13.667 , and Chamfer/Hausdorff overlap against the encoded attractor.

3.2 Pluggable predictor in a DreamerV3-style agent (Part 2)

The treatment predictor is the Part-1 E-arm embedded in the RSSM:

$$z' = z + dt \left(\underbrace{J \nabla_z H(z, h)}_{\text{flow}} - \underbrace{R(z, h) \nabla_z S(z, h)}_{\text{dissipation}} \right), \quad R = LL^\top \succeq 0, \quad (1)$$

with constant canonical J (dimension 64), L initialized at 0.05, dt learnable, and H, S MLPs conditioned on the GRU memory h . The prior is Gaussian with this mean and fixed variance, keeping the KL analytic with free bits. The baseline is DreamerV3’s categorical RSSM (32 classes $\times$ 32 codes). All other components — encoder, decoder, reward/continue heads, TD(λ) imagination, actor-critic — are byte-identical across arms; any difference is attributable to the predictor structure alone.

4 Experiments

4.1 Part 1: failure atlas (3 seeds, real numbers)

Table 1 summarizes 3-seed runs (full traces in **results/**, runner in **benchmarks/**).

arm	motion ratio	model λ_1	contraction	$\text{tr}(R)/n$
B unconstrained	—	explodes	—	—
C rigid Hamiltonian	~ 0	~ 0	≈ 0	—
D naive metriplectic	low	overshoot	wrong	0.00000
E fixed metriplectic	1.1–2.9	1.0–4.3	$-\mathbf{13.80} \pm \mathbf{0.17}$	2.3–6.2
F + spectral align	1.2–3.8	1.4–5.0	-13.25 ± 0.19	2.5–8.3

Table 1: Failure atlas on lifted Lorenz-63. True divergence -13.667 ; true $\lambda_1 \approx 0.9$. E recovers contraction to $\sim 1\%$, but no arm pins λ_1 ; the rigid arm posts the best *overlap* while being dynamically dead.

Reading. (i) The dissipation gate works: once J is constant and R escapes the dead saddle, learned contraction matches truth. (ii) Spectral shaping does not: pinning the *sum* of the spectrum cannot pin any individual Oseledets exponent. (iii) Overlap metrics reward dead models; they must be gated on rollout liveness.

4.2 Part 2: Crafter sample efficiency (controlled)

Protocol. Two arms (RSSM baseline, metriplectic treatment), identical seeds and budgets; 20k environment steps; world model trained on batches of 16×16 -frame sequences; evaluation every 5k steps by greedy rollout return. KL free bits use a matched *total* budget of 32 bits across arms (1.0 per categorical code, as in DreamerV3, vs. 0.5 per continuous dimension), so posterior-collapse slack cannot confound the comparison. Every 2k steps the agent logs a latent-stability audit: raw (uncapped) prior/posterior KL, prior-vs-posterior drift normalized by latent scale, latent boundedness, and — for the metriplectic arm — dissipation engagement $\text{tr}(R)/n$. Runs execute with a hard wall-clock budget on a T4 (exact commands: `experiments/crafter/train.py`, `kaggle/README.md`).

Hypotheses. H_+ : structure helps — the metriplectic arm needs fewer environment steps to reach a given return, because its prior mean is a better dynamics prior (less posterior collapse, longer useful imagination). H_- : structure hurts — the continuous latent and learned potentials slow the world model relative to the discrete categorical prior. Either outcome is a publishable, controlled answer.

Latent stability, defined. The quantity a world model really sells is its *prior* $p(z_t | h_t)$: at imagination time the actor rolls out the prior alone, and any gap between what the prior predicts and where the posterior lands after the true observation is error the agent acts on. We therefore instrument, every 2k steps and out-of-distribution from the loss (the audit runs under `no_grad` on a fresh batch): (i) *raw KL* $D_{\text{KL}}(q \| p)$ before free-bit clamping — the predictive error the loss actively hides; (ii) *prior-posterior drift* $\mathbb{E}\|z_t^p - z_t^q\|^2$, reported scale-free as $\text{drift}/\|z\|^2$ so the 1024-D one-hot categorical latent and the 64-D continuous latent are comparable; (iii) *latent boundedness* $\mathbb{E}\|z\|^2$; and (iv) for the treatment arm, *dissipation engagement* $\text{tr}(R)/n$ — does the metriplectic map actually dissipate inside the RL loop, or does it sit on the dead saddle again? The hypothesis pairing: if structure helps, the metriplectic prior’s drift and raw KL should stay below the RSSM’s at matched environment steps (better dynamics prior $\Rightarrow$ longer useful imagination); if structure hurts, the continuous

arm	return @ 20k	drift/ $\ z\ ^2$	raw KL
RSSM (DreamerV3)	–	–	–
Metriplectic (E-arm)	–	–	–

Table 2: Crafter sample efficiency and latent stability (3 seeds, mean $\pm$ std), to be filled from the T4 runs (see [kaggle/README.md](#)). Placeholders are deliberate: we do not ship numbers we have not run.

latent’s drift should exceed the categorical’s despite matched free-bit budgets.

Preliminary verification (CPU, fake env). Both arms train end-to-end from a clean checkout: reconstruction loss falls, KL free bits active, actor/critic losses finite, evaluation episodes complete, and the stability audit runs — the metriplectic arm’s dissipation engages ($\text{tr}(R)/n > 0$ and rising) inside the RL loop, mirroring the Lorenz atlas. This validates the pipeline; it is not evidence about the hypothesis. The headline comparison (Table 2 and the stability curves) is produced by the shipped T4 runs.

5 Related work

Structure-inducing predictors: Hamiltonian neural networks [6, 4, 16], metriplectic/GENERIC learning [17, 11, 7], latent-world-model physics [1, 13]. JEPA and model-based RL: [12, 2, 5, 9, 10, 3]. This paper differs in two ways: it documents *which* structural failure mode each class exhibits (the atlas), and it runs the first controlled predictor-ablation inside a DreamerV3-style agent on the benchmark (Crafter) of the group that introduced RSSM.

6 Conclusion

Physical structure in latent predictors is not a free lunch: each naive class fails on a specific, diagnosable property. The fixes that make dissipation genuinely engage (constant divergence-free J ; broken dead saddle) are concrete and validated. Whether those fixes pay off inside a real RL world model is the open question we answer with the Crafter experiment — shipped, reproducible, and bounded by an hour of compute. The code, runs, and drafts are public; we invite scrutiny of every number.

References

- [1] K. Allen, K. Smith, J. Tenenbaum. Latent world models for intuitive physics. *Trends in Cognitive Sciences*, 2022.
- [2] M. Assran, Q. Duval, I. Misra, P. Bojanowski, P. Vincent, M. Rabbat, Y. LeCun, N. Balas. Self-supervised learning from images with a joint-embedding predictive architecture. CVPR 2023.

- [3] J. Bruce, M. Dennis, A. Edwards, J. Parker-Holder, Y. Shi, E. Hughes, M. Lai, A. Mavalankar, R. Steigert, S. Koyejo, P. Kumar, R. Gosula, O. Farquhar, T. Mesnard, T. Ritchie, S. Gheshlaghi Azar, M. Valko, N. Heess, T. Ewalds, R. Hadsell, L. Espeholt. Genie: Generative interactive environments. ICML 2024.
- [4] Z. Chen, J. Zhang, M. Arjovsky, L. Bottou. Symplectic recurrent neural networks. ICLR 2020.
- [5] Q. Garrido, M. Assran, M. Ballas, A. Bardes, L. Najman, Y. LeCun. Learning and leveraging world models in visual representation learning. 2024.
- [6] S. Greydanus, M. Kater, J. Kaplan, B. Kakade. Hamiltonian neural networks. NeurIPS 2019.
- [7] S. Ha, H. Jeong. Dissipative symODEN. ICLR 2023.
- [8] D. Hafner. Benchmarking sparse reward environments with Crafter. *arXiv:2109.08714*, 2021.
- [9] D. Hafner, J. Pasukonis, J. Ba, T. Lillicrap. Mastering diverse domains through world models. *arXiv:2301.04104*, 2023.
- [10] D. Hafner, T. Lillicrap, M. Norouzi, J. Ba. Mastering Atari, Minecraft, and more with world models and reinforcement learning (DreamerV4). 2024.
- [11] Q. Hernandez, A. Badias, D. Gonzalez, F. Chinesta, E. Cueto. Deep learning of thermodynamics-aware reduced-order models from data. *Phys. Rev. E*, 2021.
- [12] Y. LeCun. A path towards autonomous machine intelligence, 2022.
- [13] K. Martynov, A. Mehrotra, A. Rajkumar. Structure-preserving latent world models. 2023.
- [14] V. Oseledets. A multiplicative ergodic theorem. *Trans. Moscow Math. Soc.*, 1968.
- [15] H. Ottinger. *Beyond Equilibrium Thermodynamics*. Wiley, 1997.
- [16] P. Toth, D. Rezende, A. Jaegle, S. Racaniere, A. Botev, I. Higgins. Hamiltonian generative networks. ICLR 2020.
- [17] Z. Zhang, S. Shin, G. Karniadakis. GFINNs: GENERIC formalism informed neural networks. *CMAME*, 2021.